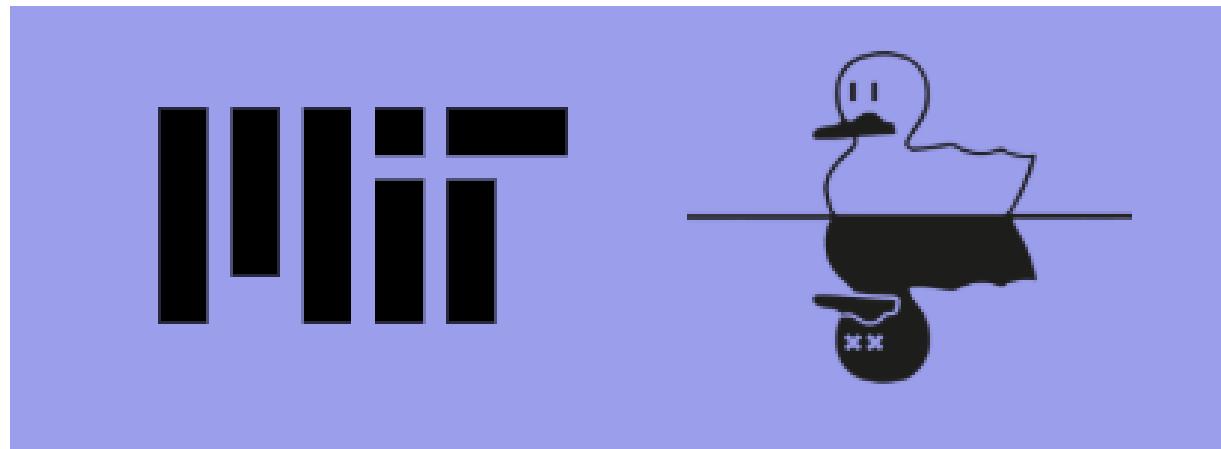


# FluxQuapacitors vs LABS: A CUDA-Q Acceleration Story



NVIDIA®



o.qBraid

# TEAM MEMBER DETAILS

**Project Lead (Architect): Vedika Saravanan**

**GPU Acceleration PIC (Builder): Lim Wee Keat and Gabriel Ortega**

**Quality Assurance PIC (Verifier): Gabriel Ortega and Yash Singh**

**Technical Marketing PIC (Storyteller):Travis Martin**

# The LABS Problem

**Low Autocorrelation Binary Sequence (LABS) is a fundamental problem in combinatorial optimization and signal processing.**

- **Goal:** Find a binary sequence of length  $N$  that minimize aperiodic autocorrelation side-lobes → maximize merit factor / minimize energy
- **Scale:** Exponential search space of  $2^n$  configurations
- **Difficulty:** Highly rugged, non-convex energy landscape → classical heuristics often get trapped in local minima
- **Why it matters:** High-precision radar, deep-space pulse compression, and cryptographic keys

→ LABS is therefore an ideal candidate for hybrid quantum-classical acceleration with **CUDA-Q**.

# LABS as an Energy Minimization Problem

## TOTAL ENERGY

$$E = \sum_{k=1}^{N-1} C_k^2$$

3

Lower energy indicates better autocorrelation properties.

Minimum Energy (Example, N = 7)

This energy value defines sequence quality – reaching the global minimum directly enables high-precision signal detection in hardware.

# **Plan and Pivot**

# Original PRD strat:

**Objective:** Scale and refine the Quantum-Enhanced Memetic Tabu Search (QE-MTS) for larger LABS instances.

- **QAOA optimization**
  - Replace fixed angles with adaptive optimization
  - Use COBYLA / SPSA to tune variational parameters ( $\gamma$ ,  $\beta$ )
  - Improves quality of the initial candidate pool
- **GPU-accelerated simulation**
  - Migrate backend to NVIDIA cuQuantum SDK
  - Enables deeper circuits and larger qubit counts
- **Enhanced classical search**
  - Add adaptive tabu tenure and diversification
  - Prevents premature convergence in rugged, high-N landscapes

→ **Goal: Robust performance as LABS scales to larger N**

# What changed?

## QAOA parameter strategy

- PRD: Fully variational QAOA loop using COBYLA / SPSA
- Final: Fixed ansatz parameters
- Why: Prioritized rapid sampling and clean hardware benchmarking isolation

## Tabu Search configuration

- PRD: Adaptive tabu tenure and diversification
- Final: Fixed-tenure tabu search
- Why: Reduced classical overhead within the hackathon timeframe

## CUDA-Q kernel constraints

- Unplanned change: Manual decomposition of R(ZZ)
- Implementation: CNOT → RZ → CNOT
- Reason: Ensure compatibility with the available CUDA-Q gate set

→ **Outcome: Faster iteration, stable benchmarking, and a fully CUDA-Q-compatible pipeline**

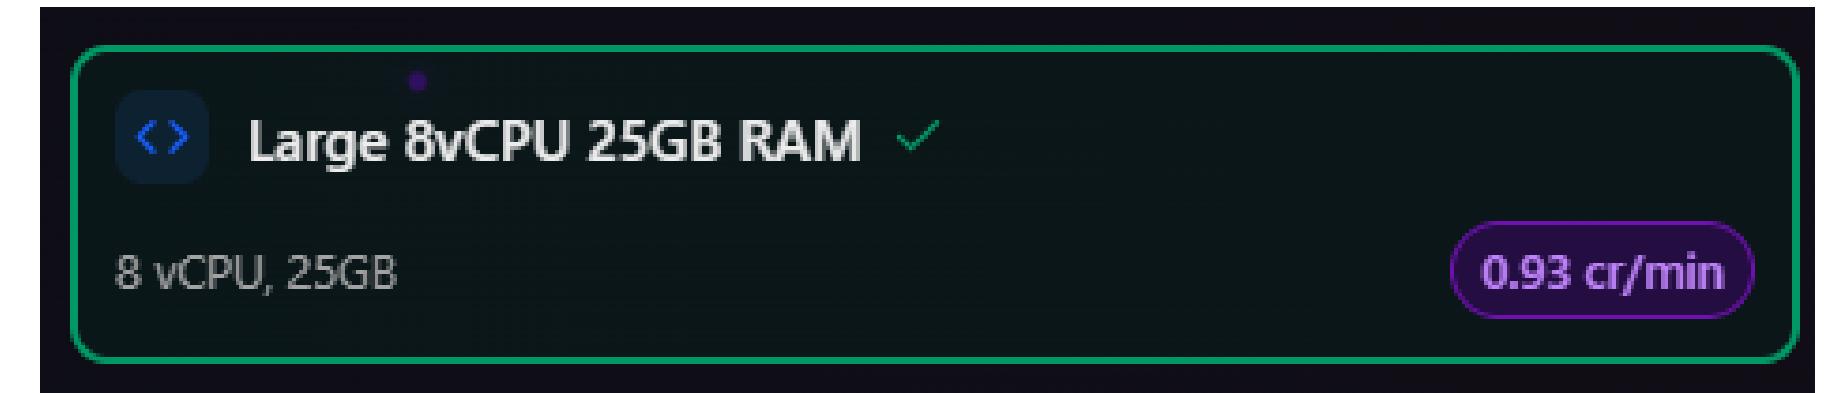
# **Results!!!**

# CPU vs GPU Benchmarking Setup

**Goal:** Compare solution quality and runtime across CPU and GPU backends under identical conditions.

## Hardware compared

- CPU: qBraid Large 8 vCPU (25 GB RAM)
- GPUs: NVIDIA Brev GPUs
- Control: Local RTX 4080 (physical GPU)



## Experimental controls

- Same algorithm, parameters, and problem instances
- Solution graph held constant
- Only variable: hardware backend

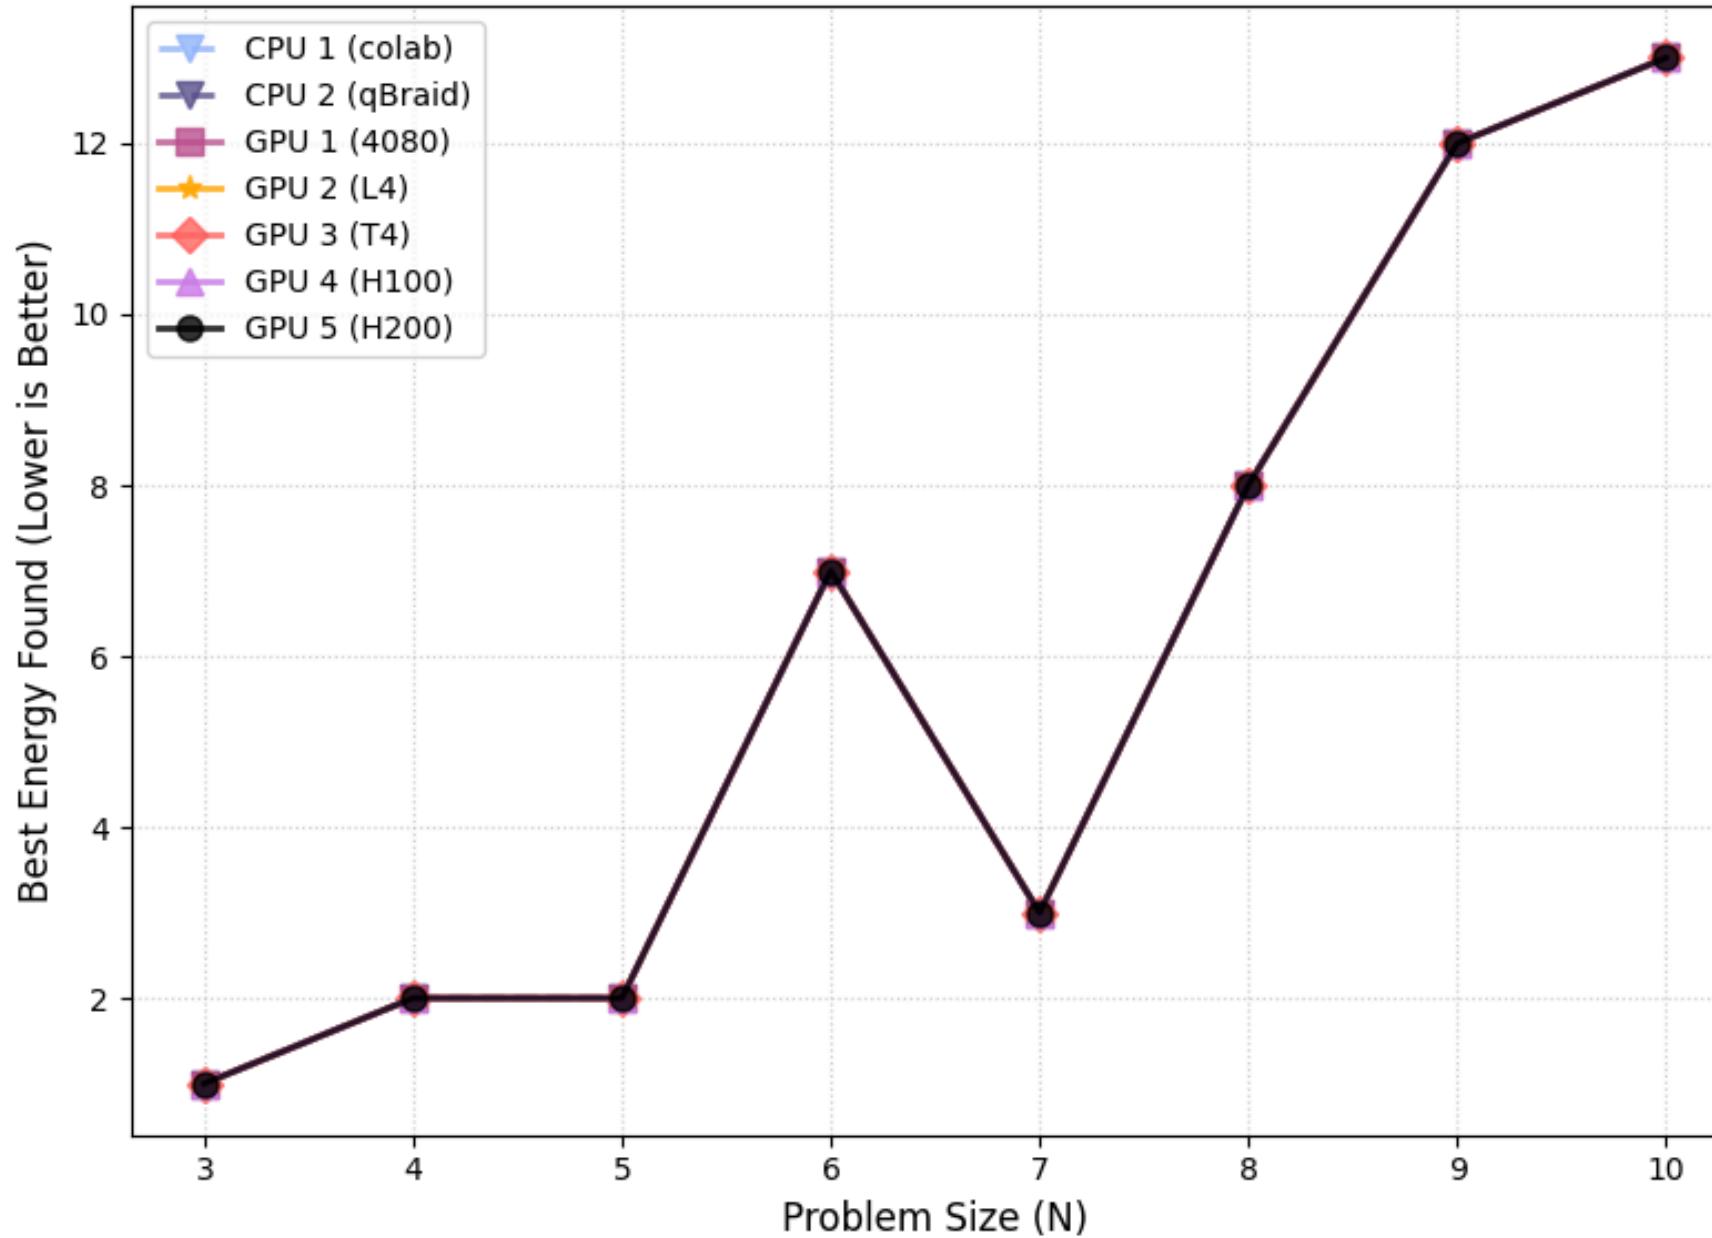
## Why this matters

- Isolates true hardware effects
- Ensures fair, apples-to-apples comparison
- No algorithmic shortcuts or hidden optimizations

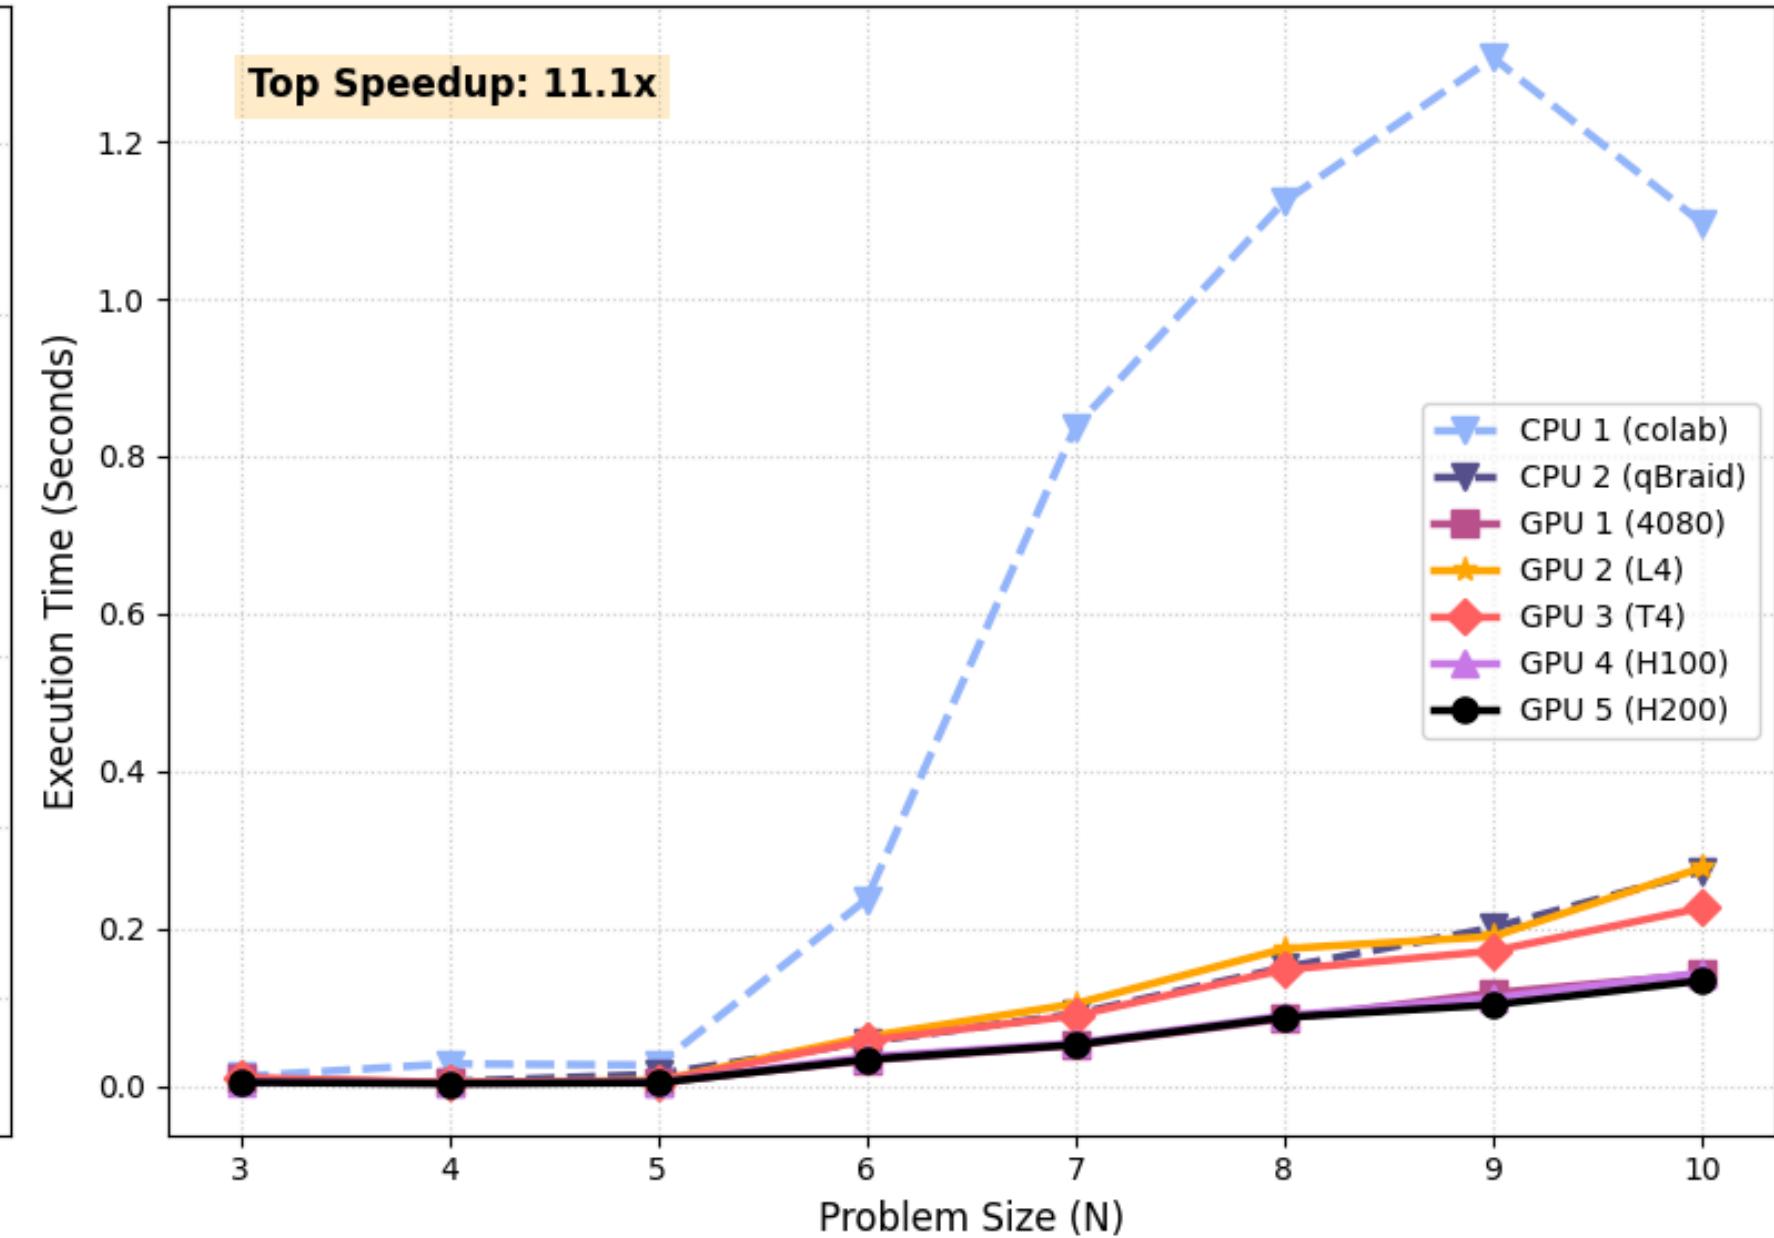
→ **Outcome: Differences in performance reflect hardware capabilities only**

## Hardware Benchmark: Hybrid Quantum Optimization (QE-MTS)

Optimization Quality (Energy)

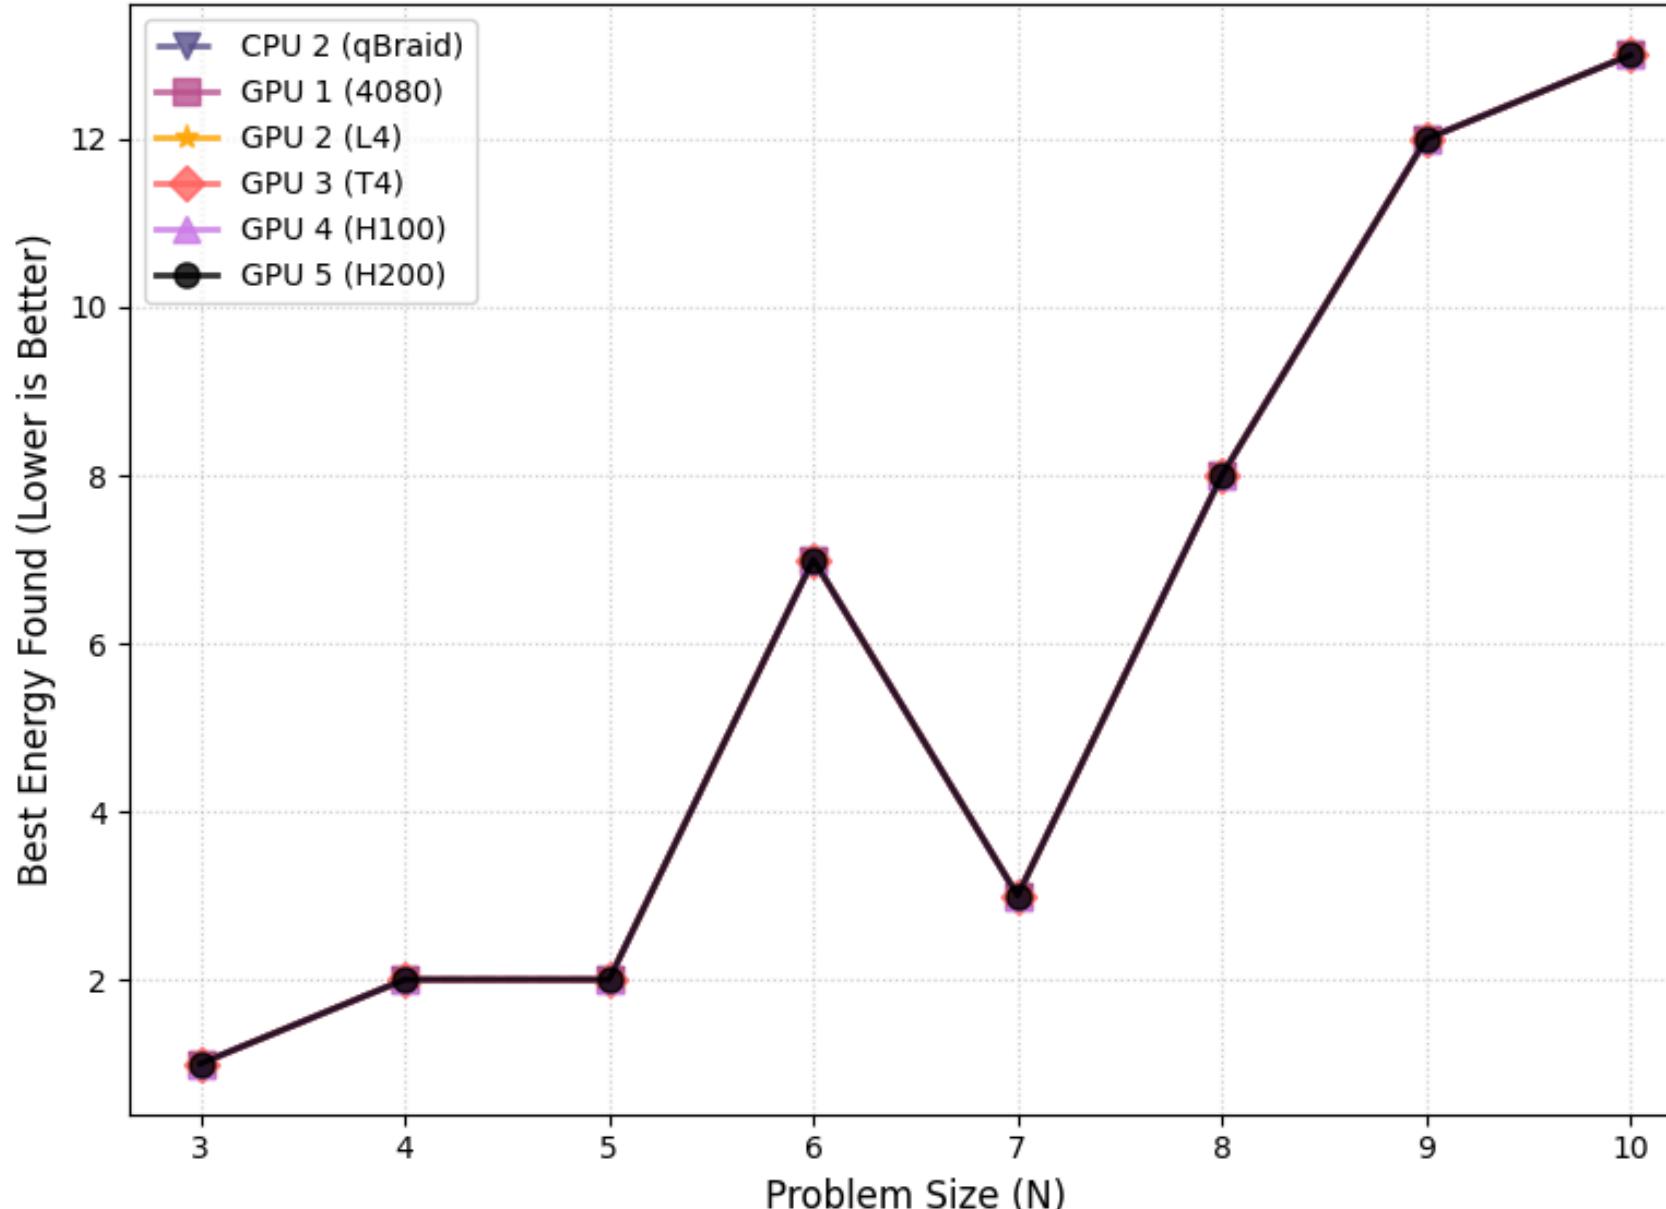


Runtime Performance (Speed)

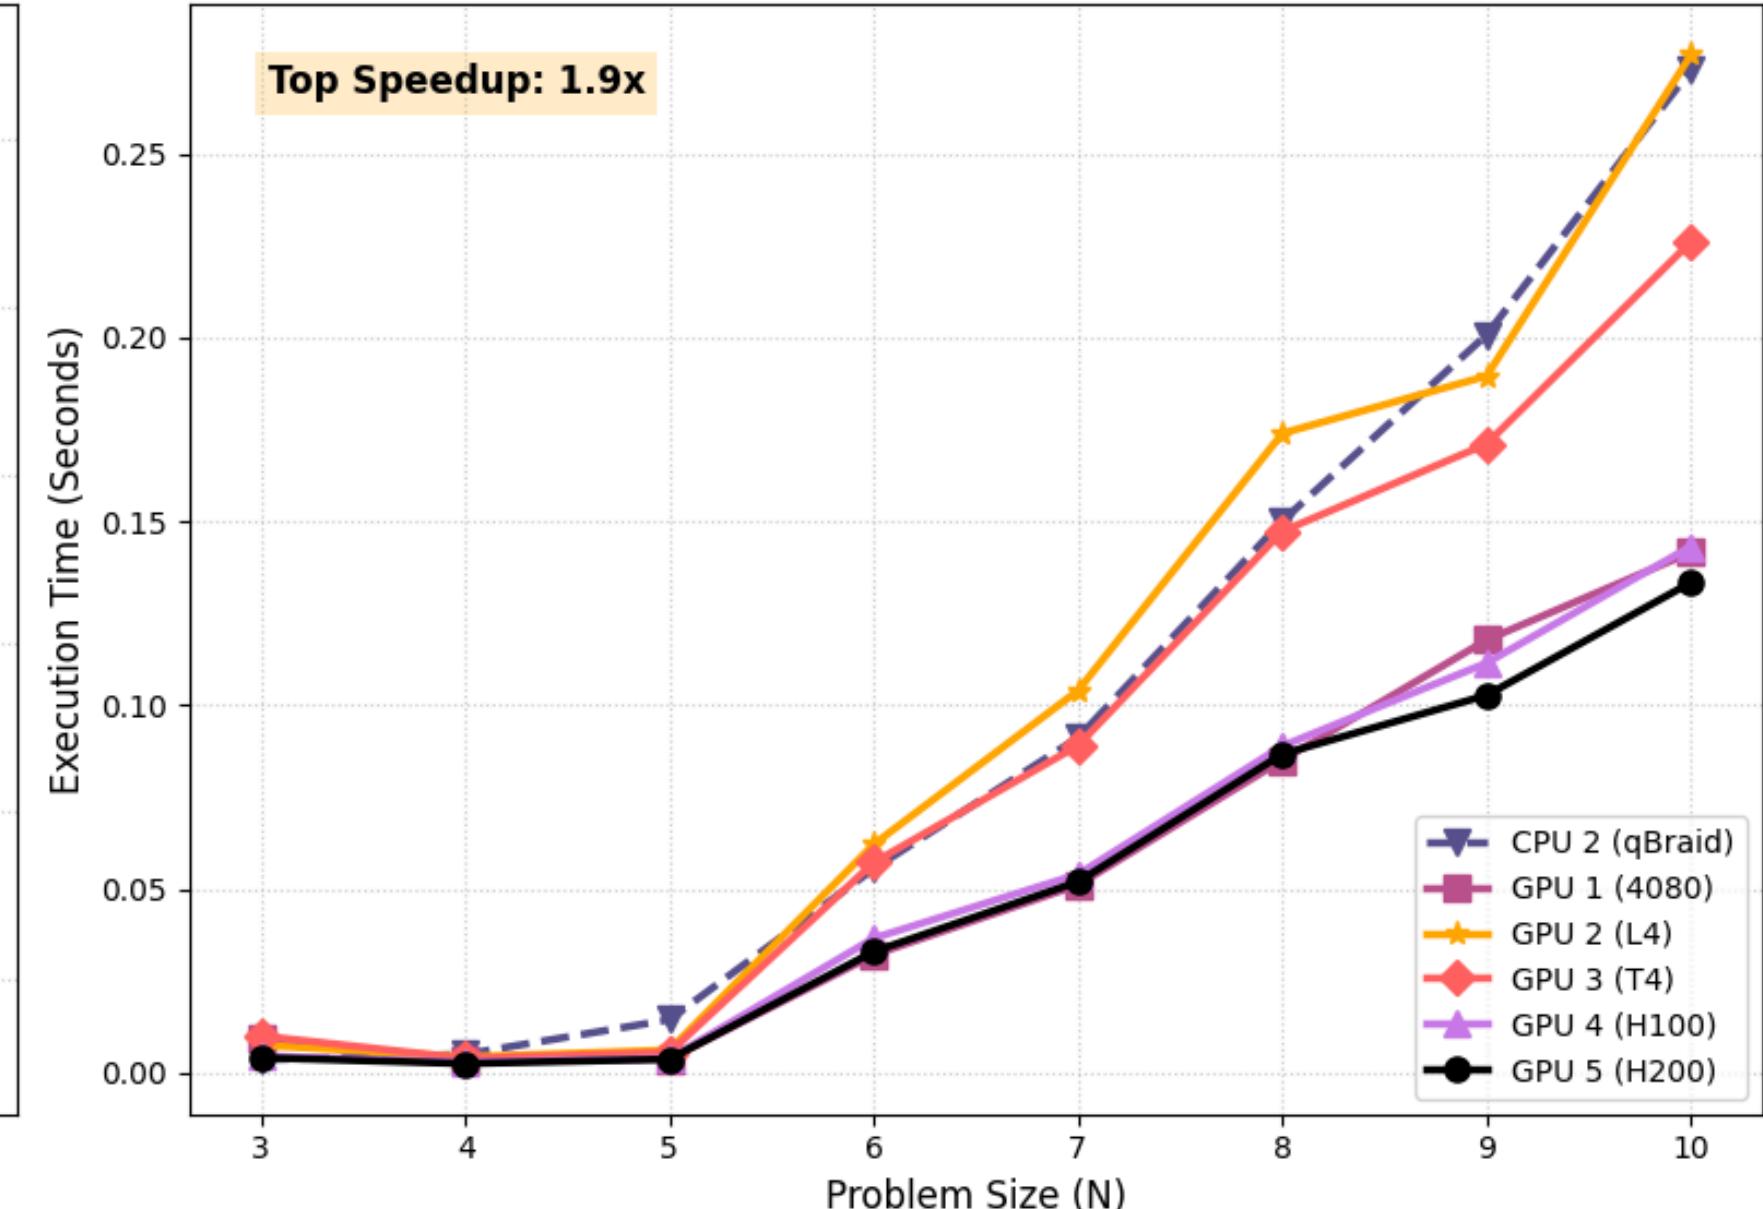


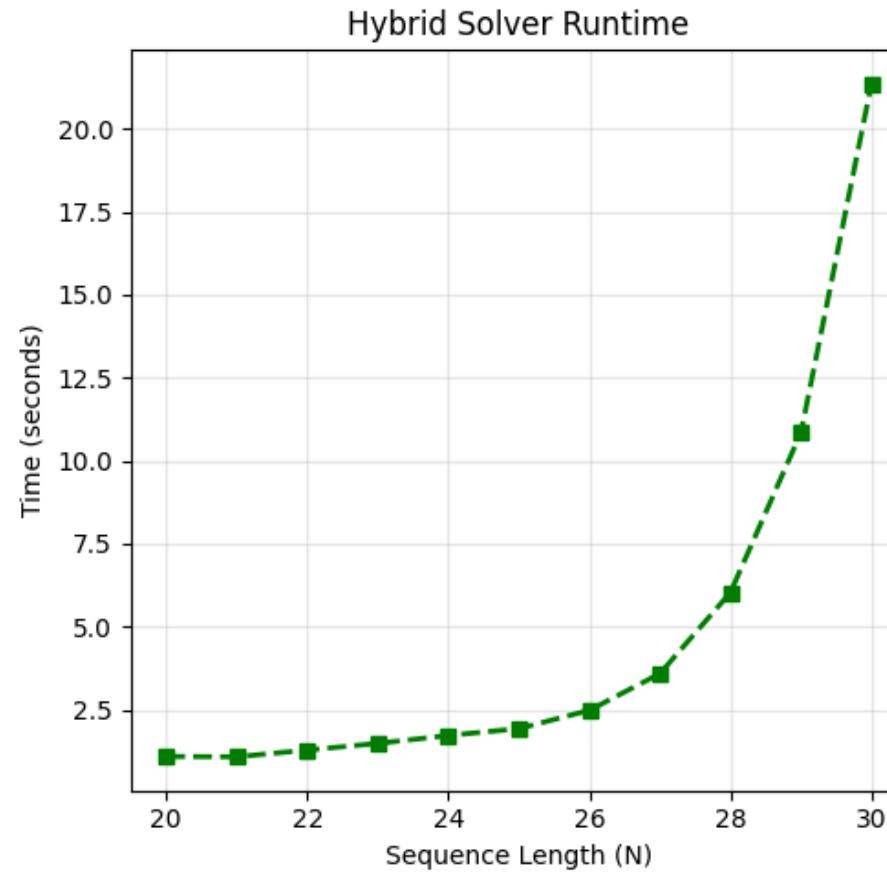
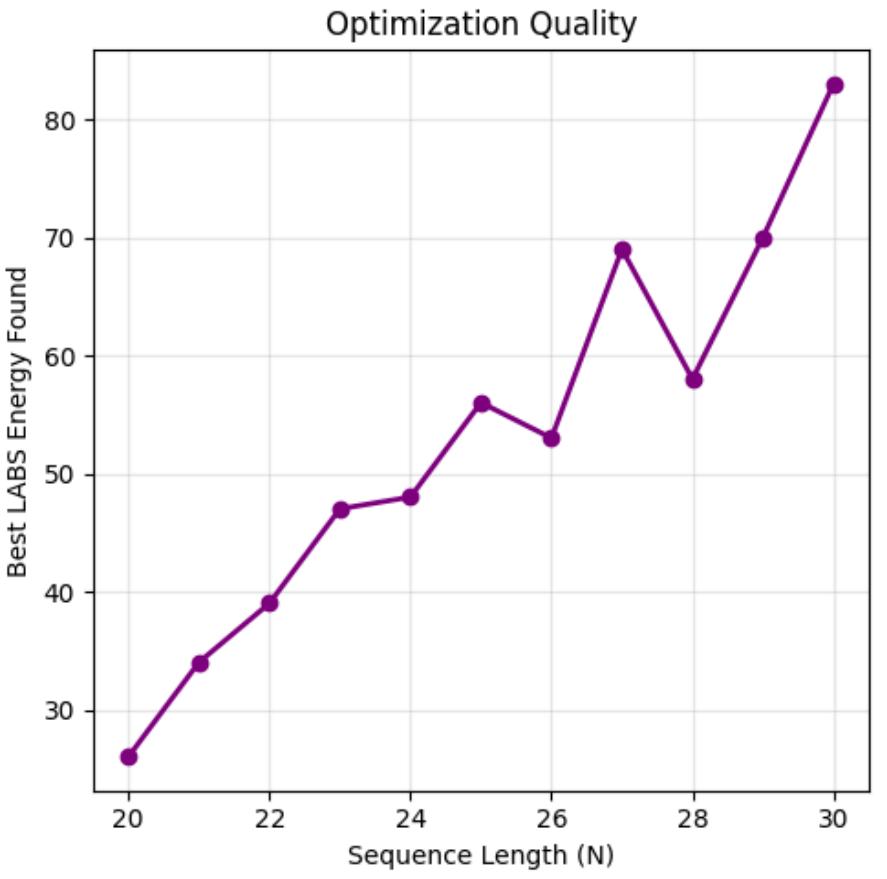
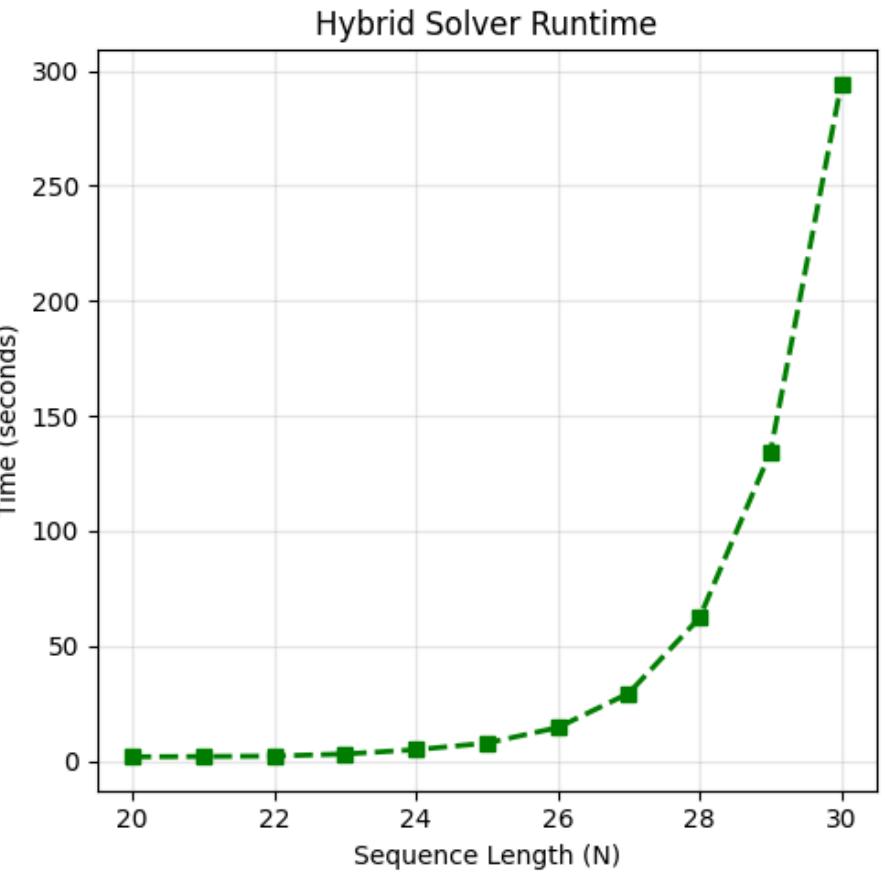
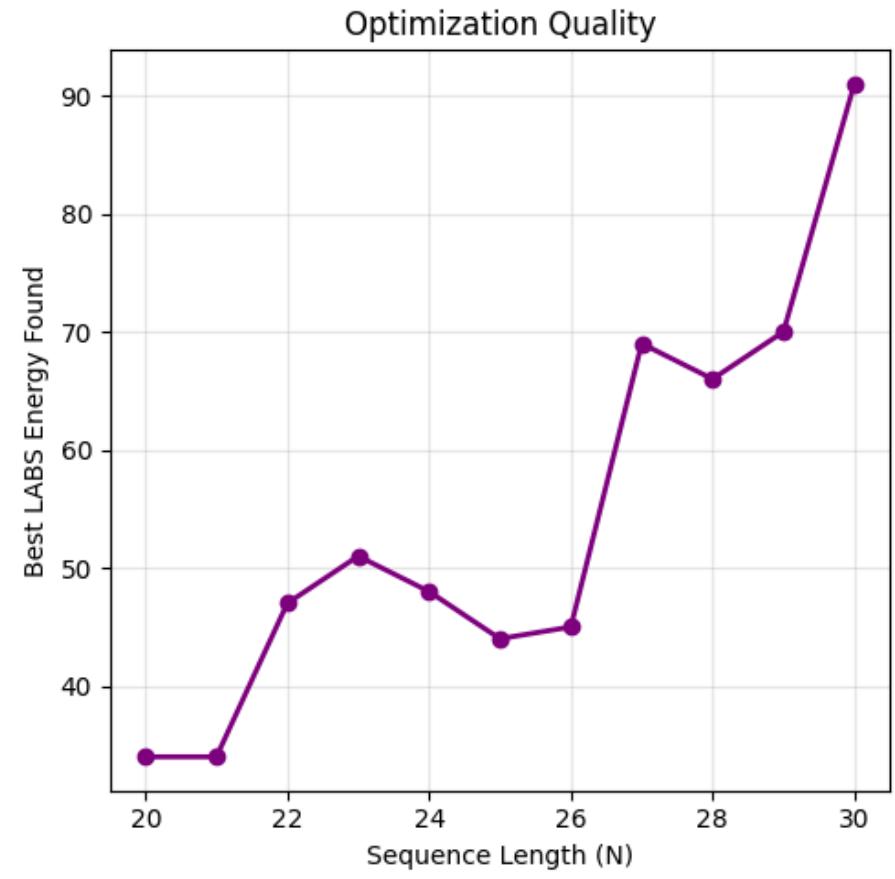
## Hardware Benchmark: Hybrid Quantum Optimization (QE-MTS)

Optimization Quality (Energy)

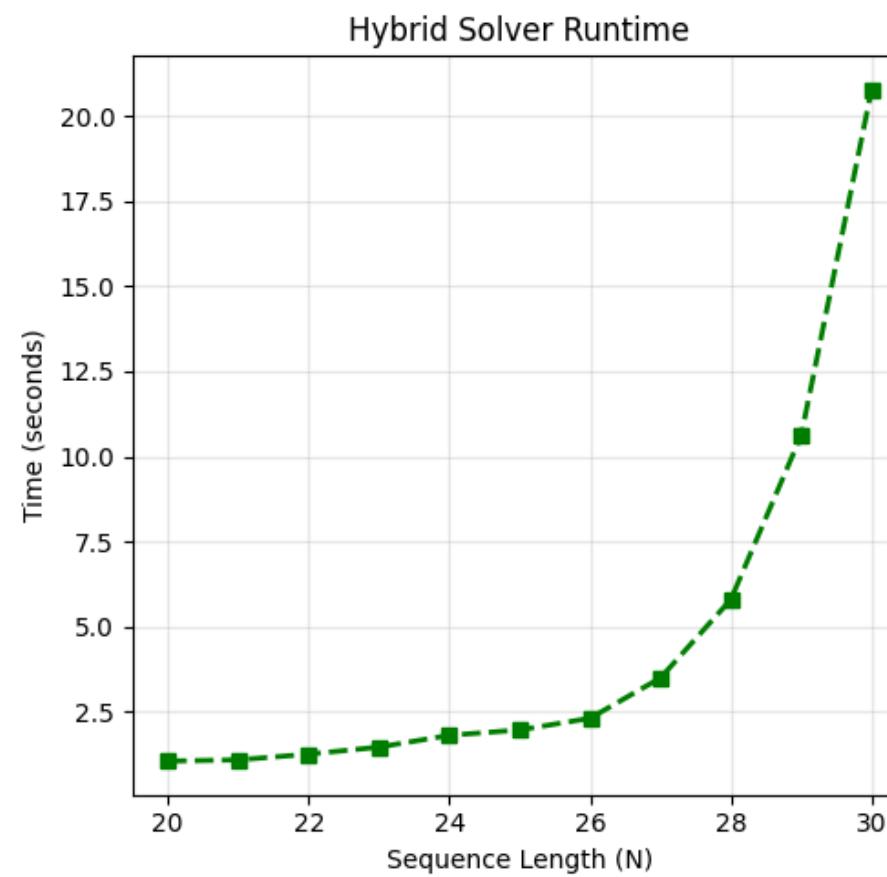
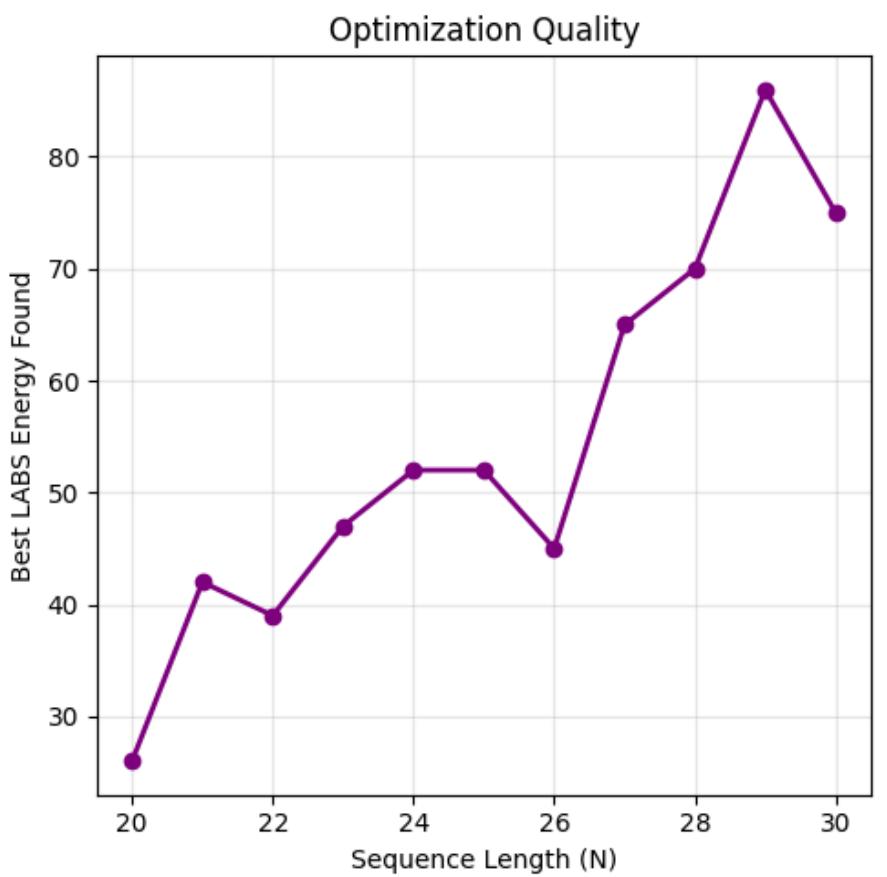
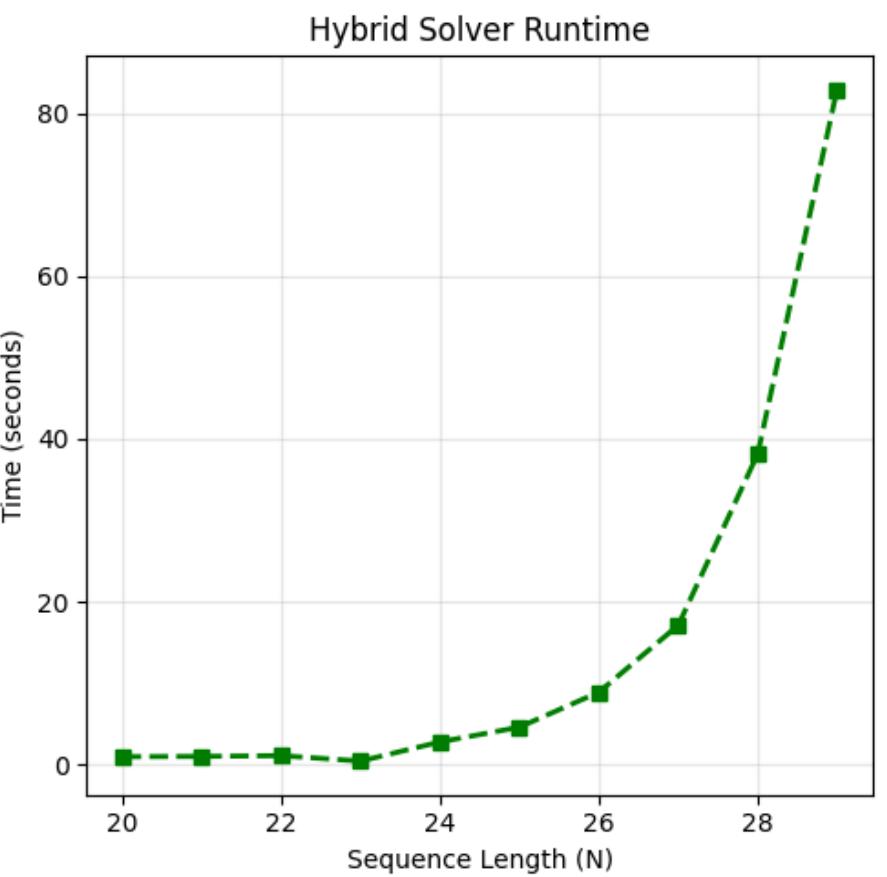
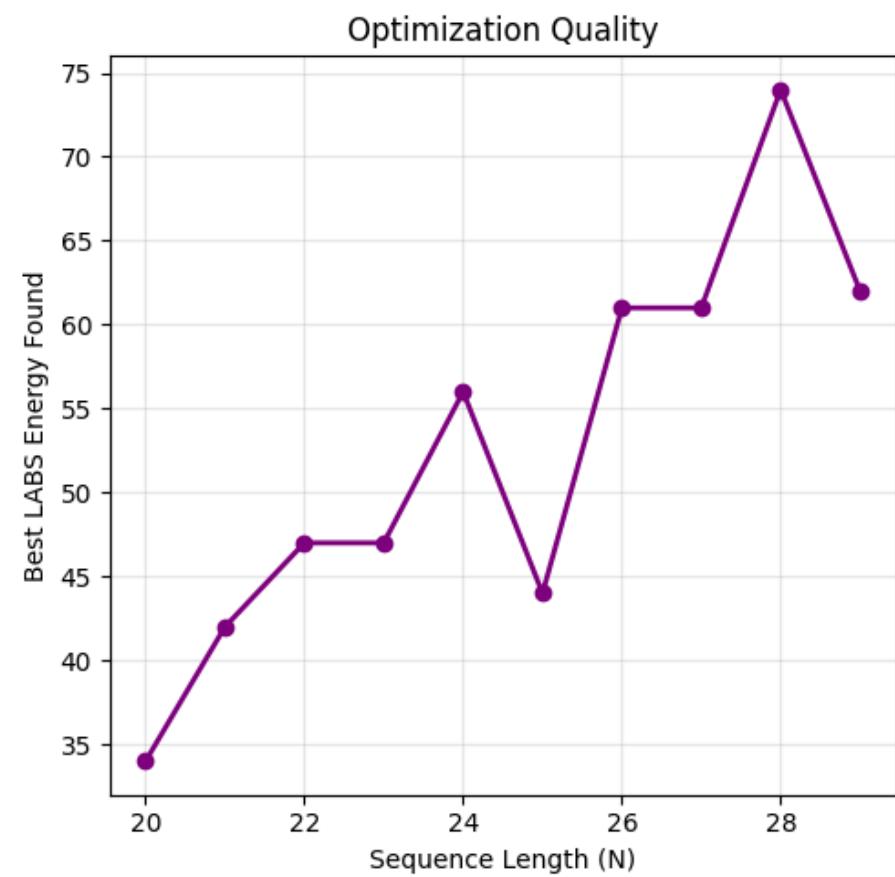


Runtime Performance (Speed)





**L4**



**RTX 4080**

**H200**

# Summary

## Key Observations & Takeaways

- Algorithm quality alone does not determine performance at scale
- Hardware characteristics become dominant as problem size increases
- **Small scale ( $n = 3 \rightarrow 10$ ):**
  - No meaningful difference between CPU and GPU results
  - Output is stable across hardware types
- **Larger scale ( $n = 20 \rightarrow 30$ ):**
  - Significant variation in outputs for the same algorithm
  - Differences driven primarily by GPU quality and architecture
- **Runtime performance is not a reliable predictor of solution quality**
- **Example:**
  - RTX 4080 vs H200 show similar runtimes
  - Yet produce different energy outputs
- **Implication:**
  - GPU choice impacts numerical outcomes, not just speed
  - Hardware-aware evaluation is critical for large-scale workloads